

SANS

OFFENSIVE
OPERATIONS

HACKFEST

SUMMIT 2023



November 16 - 17

Hollywood, CA & FREE Live Online



Just Pocket Change? Examining The Cost Of “Nickel and Dime” Exploit Mitigations

Connor McGarr

HACKFEST
SUMMIT 2023



About Me

- Software Engineer @ CrowdStrike
 - EDR features, detection content, etc.
- Blog
 - <https://connormcgarr.github.io>
- I like low-level Windows stuff!



What Are “Nickel and Dime” Mitigations?

- We all have that “spare change” jar
 - ...well maybe you don’t, but my family does
- Although the spare change jar is full of coins that, by themselves, are seemingly insignificant one day you wake up and realize you have quite a lot of money in the jar!



What Are “Nickel and Dime” Mitigations?

- Enter “nickel and dime” mitigations
 - These are the “smaller” mitigations we may not give a second look
 - These mitigations also may not mean a lot “by themselves”
 - But when combined they can significantly raise the cost of exploitation, just like coins in a spare change jar!



Exploitation Trends

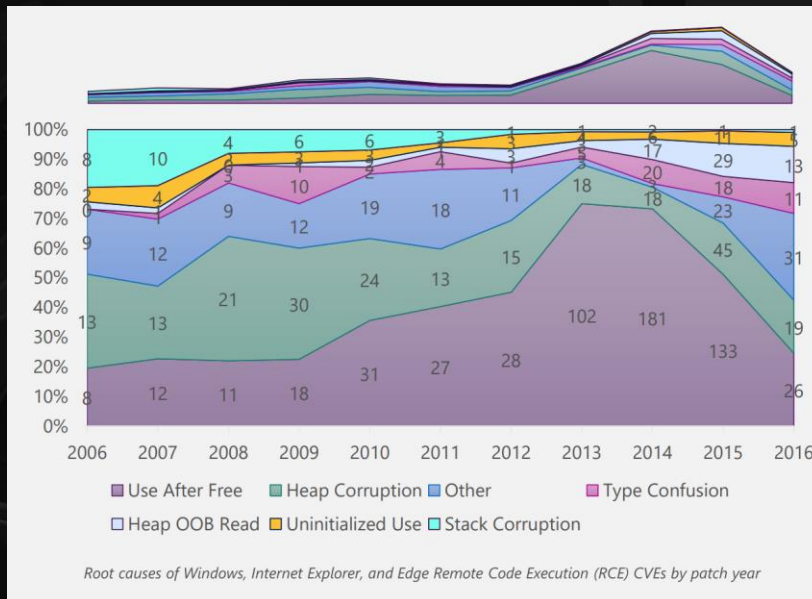
- Before we get into the specific mitigations themselves, it's worth considering how exploitation has changed over time
- Exploit mitigations have been around for awhile
 - DEP/ASLR have been supported since Windows XP/Vista
 - Control Flow Integrity (CFG in the case of Microsoft) since Windows 8.1 as an optional update

Exploitation Trends

- The origin of mitigations...
 - Most exploits in “old” days are usually (but not always) stack overflows or use-after-free vulnerabilities
 - Both require making memory executable (mitigated by DEP)
 - Both require knowledge of the address space (mitigated by ASLR)
 - Use-after-frees (in many cases) specifically require crafting a fake object to overwrite a function pointer (mitigated by CFI)

Exploitation Trends

- DEP/ASLR/CFI, outlined by Microsoft in 2016, made a significant impact regarding stack corruption and use-after-free vulnerabilities
 - This success was used as an opportunity to invest heavily into “modern” mitigations around this time



Stack corruption issues have been essentially eliminated

Use after free issues rose dramatically in 2013 & 2014 but have since decreased

Heap out-of-bounds read, type confusion, and DLL planting have increased

Exploitation Trends

- This investment led to several well-known mitigations:
 - **Arbitrary Code Guard (ACG)**
 - Memory is either RW or RX, but never both
 - **Code Integrity Guard**
 - Only Microsoft signed DLLs can be loaded
 - **Intel CET and eXtended Control Flow Guard (XFG)**
 - Hijacking control flow is essentially impossible/very limited
 - **Supervisor Mode Execution (and Access) Prevention (SMEP/SMAP)**
 - Prevents redirecting execution/data access into user-mode, from kernel-mode (with some caveats not mentioned here)
 - **Hypervisor Protected Code Integrity (HVCI)**
 - Memory permissions are managed by hypervisor, can't create RWX memory dynamically in kernel-mode (the "ACG" of kernel-mode)
- ...and many more not mentioned!

Exploitation Trends

- Even with the heavy investment in mitigations, we know exploitation is still occurring. Why?
 1. Mitigations may still be disabled
 2. There are still other issues that need to be addressed
 - Generic memory leaks
 - Full implementation of things like ASLR in kernel-mode
 - Data-only corruption techniques
 - Other techniques which don't rely on executable code creation or control-flow hijacking
- Today we will look at how some of the “smaller”, but often overlooked, mitigations affect exploitation cost and how they directly address some of these lingering issues

“Nickel and Dime” Mitigations

- Full ASLR Example 1: Page Table Randomization
 - Page tables maintain entries known as page table entries (PTEs)
 - PTEs are used to determine what physical frames correspond to what physical addresses, and to enforce virtual memory permissions (readable/writable/etc.)

```
Command X
0: kd> dx @cr3
@cr3 : 0x7d4000 [Type: unsigned __int64]
0: kd> !vtop 0x7d4000 0xffffffff7800000000
Amd64VtoP: Virt fffff78000000000, pagedir 00000000007d4000
Amd64VtoP: PML4E 00000000007d4f78
Amd64VtoP: PDPE 0000000004500000
Amd64VtoP: PDE 0000000004301000
Amd64VtoP: PTE 0000000004302000
Amd64VtoP: Mapped phys 000000000182b000
Virtual address fffff78000000000 translates to physical address 182b000.
```

“Nickel and Dime” Mitigations

- Full ASLR Example 1: Page Table Randomization
 - Page tables reside in physical memory
 - However, CPU instructions can only directly access virtual memory
 - This is a problem as paging structures need to constantly be updated/flushed/etc.
 - How can we access them if we are limited to accessing virtual memory only?
 - To compensate, the OS maps all the PTEs for each of the paging structures into virtual memory in the form of a “per-structure array”

```
Command X
0: kd> !pte 0xfffff78000000000
VA fffff78000000000
PXE at FFFFF8FC7E3F1F78 PPE at FFFFF8FC7E3EF000 PDE at FFFFF8FC7DE00000 PTE at FFFFF8FBC0000000
contains 000000004500063 contains 000000004301063 contains 000000004302063 contains 800000000182B963
pfn 4500 ---DA--KWEV pfn 4301 ---DA--KWEV pfn 4302 ---DA--KWEV pfn 182b -G-DA--KW-V
```

- Because each of these PTE arrays is needed early in the OS loading processes, they were stored at a static address


| Address | Function | Instruction |
|-----------------------|--------------|-----------------------------|
| INIT:0000000140BC3F79 | MiInitSystem | mov rax, 0FFFFFF6800000000h |

Static PTE array for 4KB-aligned pages

“Nickel and Dime” Mitigations

- Why is this bad?
 - An adversary, **without** a read primitive, can correlate a memory address to its PTE
 1. Obtain the virtual page number (VPN) by dividing the virtual address by sizeof(PAGE_SIZE)
 2. Obtain the index into the array (multiply the VPN by sizeof(PTE))
 3. Index the array (**add the index to the address of the PTE array**)
 - With the address of the PTE, and a write primitive, an attacker can corrupt the PTE metadata, which enforces virtual memory permissions (RWX, RW, RX, copy-on-write, etc.)

nt!MiGetPteAddress
performs this with compiler
optimizations



```
; __int64 __fastcall MiGetPteAddress(unsigned __int64)
MiGetPteAddress proc near
shr     rcx, 9
mov     rax, 7FFFFFFFF8h
and     rcx, rax
mov     rax, 0FFFFFF680000000h
add     rax, rcx
retn
MiGetPteAddress endp
```

“Nickel and Dime” Mitigations

- With the array being randomized, the process for locating a page table entry doesn't change, but it *requires* a read primitive

1. Fetch virtual page number

2. Grab the index (VPN * 0x8) and add the PTE base

```
Command X
3: kd> u 0x7ff65e543600 L3
csrss!HotPatchSpareGlobal7+0x8:
00007ff6`5e543600 90          nop
00007ff6`5e543601 90          nop
00007ff6`5e543602 90          nop
3: kd> dx @$vpn = (0x7ff65e543600 / 0x1000)
@$vpn = (0x7ff65e543600 / 0x1000) : 34349638979
3: kd> dx @$pteAddr = (@$vpn * 0x8) + @$pteBase
@$pteAddr = (@$vpn * 0x8) + @$pteBase : -7971540096488
3: kd> dx *(nt!_MMPTE_HARDWARE*)@$pteAddr
*(nt!_MMPTE_HARDWARE*)@$pteAddr [Type: _MMPTE_HARDWARE]
[+0x000 ( 0: 0)] Valid : 0x1 [Type: unsigned __int64]
[+0x000 ( 1: 1)] Dirty1 : 0x1 [Type: unsigned __int64]
[+0x000 ( 2: 2)] Owner : 0x1 [Type: unsigned __int64]
[+0x000 ( 3: 3)] WriteThrough : 0x0 [Type: unsigned __int64]
[+0x000 ( 4: 4)] CacheDisable : 0x0 [Type: unsigned __int64]
[+0x000 ( 5: 5)] Accessed : 0x1 [Type: unsigned __int64]
[+0x000 ( 6: 6)] Dirty : 0x1 [Type: unsigned __int64]
[+0x000 ( 7: 7)] LargePage : 0x0 [Type: unsigned __int64]
[+0x000 ( 8: 8)] Global : 0x0 [Type: unsigned __int64]
[+0x000 ( 9: 9)] CopyOnWrite : 0x0 [Type: unsigned __int64]
[+0x000 (10:10)] Unused : 0x0 [Type: unsigned __int64]
[+0x000 (11:11)] Write : 0x1 [Type: unsigned __int64]
[+0x000 (47:12)] PageFrameNumber : 0x289c8 [Type: unsigned __int64]
[+0x000 (51:48)] ReservedForHardware : 0x0 [Type: unsigned __int64]
[+0x000 (55:52)] ReservedForSoftware : 0x0 [Type: unsigned __int64]
[+0x000 (59:56)] WsleAge : 0x2 [Type: unsigned __int64]
[+0x000 (62:60)] WsleProtection : 0x0 [Type: unsigned __int64]
[+0x000 (63:63)] NoExecute : 0x1 [Type: unsigned __int64]
3: kd> !pte 0x7ff65e543600
VA 00007ff65e543600
PXE at FFFFF8FC7E3F17F8 contains 0A0000002848A867 pfn 2848a ---DA--UWEV
PPE at FFFFF8FC7E2FFEC8 contains 0A0000002848B867 pfn 2848b ---DA--UWEV
PDE at FFFFF8FC5FFD9790 contains 0A0000002848C867 pfn 2848c ---DA--UWEV
PTE at FFFFF8BFB2F2A18 contains 82000000289C8867 pfn 289c8 ---DA--UW-V
```

@\$pteBase is now randomized and needs to **dynamically** be found

Physical frames match (this is the PTE for **0x7ff65e543600** (this is the correct PTE))

“Nickel and Dime” Mitigations

- Full ASLR Example 2: **KUSER_SHARED_DATA**
 - Attackers need to write malicious code somewhere into memory
 - The Windows kernel offers a commonly-abused technique in the form a static writable cave to store shellcode with a write primitive
 - **KUSER_SHARED_DATA** is located at **0xfffff78000000000**

```
Command X
3: kd> dx *(nt!_KUSER_SHARED_DATA*)0xfffff78000000000
*(nt!_KUSER_SHARED_DATA*)0xfffff78000000000 [Type: _KUSER_SHARED_DATA]
[+0x000] TickCountLowDeprecated : 0x0 [Type: unsigned long]
[+0x004] TickCountMultiplier : 0xfa0000 [Type: unsigned long]
[+0x008] InterruptTime [Type: _KSYSTEM_TIME]
[+0x014] SystemTime [Type: _KSYSTEM_TIME]
[+0x020] TimeZoneBias [Type: _KSYSTEM_TIME]
[+0x02c] ImageNumberLow : 0x8664 [Type: unsigned short]
[+0x02e] ImageNumberHigh : 0x8664 [Type: unsigned short]
[+0x030] NtSystemRoot : "C:\Windows" [Type: wchar_t [260]]
[+0x238] MaxStackTraceDepth : 0x0 [Type: unsigned long]
[+0x23c] CryptoExponent : 0x0 [Type: unsigned long]
[+0x240] TimeZoneId : 0x2 [Type: unsigned long]
[+0x244] LargePageMinimum : 0x200000 [Type: unsigned long]
[+0x248] AitSamplingValue : 0x0 [Type: unsigned long]
[+0x24c] AppCompatFlag : 0x0 [Type: unsigned long]
[+0x250] RNGSeedVersion : 0x9 [Type: unsigned __int64]
[+0x258] GlobalValidationRunlevel : 0x0 [Type: unsigned long]
[+0x25c] TimeZoneBiasStamp : 8 [Type: long]
[+0x260] NtBuildNumber : 0x4a62 [Type: unsigned long]
[+0x264] NtProductType : NtProductWinNt (1) [Type: _NT_PRODUCT_TYPE]
[+0x268] ProductTypeIsValid : 0x1 [Type: unsigned char]
```

“Nickel and Dime” Mitigations

- Full ASLR Example 2: **KUSER_SHARED_DATA**
 - Attackers only need a write primitive
 - This is especially true/useful for remote exploits, as KASLR is an issue (KASLR is a non-issue for local attacks, but that is changing with one of the upcoming mitigations outlined in this talk)
 - Since most memory regions/operations are page-aligned on Windows, and since **KUSER_SHARED_DATA** is ~ **0x700** bytes, that leaves ~**0x300** bytes of free memory

```
Command X
3: kd> ?? sizeof(nt!_KUSER_SHARED_DATA)
unsigned int64 0x720
```

```
3: kd> dps 0xffffffff7800000000+0x720
fffff780`00000720 00000000`00000000
fffff780`00000728 00000000`00000000
fffff780`00000730 00000000`00000000
fffff780`00000738 00000000`00000000
```

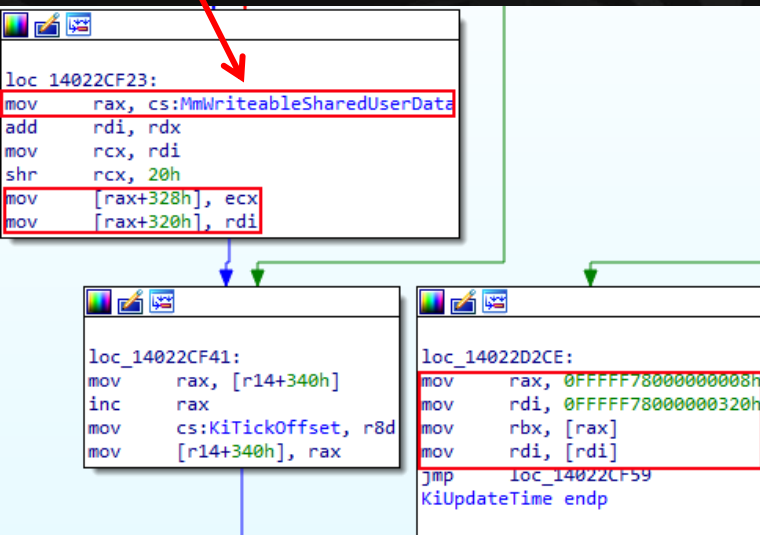
```
3: kd> !pte 0xffffffff7800000000
```

| PXE at FFFFF8FC7E3F1F78 | | PPE at FFFFF8FC7E3EF000 | | PDE at FFFFF8FC7DE00000 | | PTE at FFFFF8FBC0000000 | |
|-------------------------|-----------------|-------------------------|-----------------|-------------------------|-----------------|-------------------------|-----------------|
| contains | 000000004500063 | contains | 000000004301063 | contains | 000000004302063 | contains | 80000000182B963 |
| pfn | 4500 | pfn | 4301 | pfn | 4302 | pfn | 182b |
| ---DA--KWEV | | ---DA--KWEV | | ---DA--KWEV | | -G-DA--KW-V | |

“Nickel and Dime” Mitigations

- Full ASLR Example 2: **KUSER_SHARED_DATA**
 - In the Windows Insider Preview builds, **KUSER_SHARED_DATA** has a new randomized write-only mapping
 - For read operations, the “old” **KUSER_SHARED_DATA** is used, as it is read-only
 - For write operations to any of the fields in **KUSER_SHARED_DATA** (like updating **SystemTime**), a new mapping of **KUSER_SHARED_DATA** (**nt!MmWriteableSharedUserData**) is used

Write operation



Read operation

New mapping is backed by the same physical page. Any write to the new mapping is automatically updated in the read-only mapping

| Command | | | |
|---|---------------------------|---------------------------|--------------------------|
| 0: kd> !pte 0xfffff78000000000 | | | |
| VA fffff78000000000 | | | |
| PXE at FFFFD6E371B8F78 | PPE at FFFFD6E371EF000 | PDE at FFFFD6E3DE00000 | PTE at FFFFD67BC0000000 |
| contains 00000000089C063 | contains 00000000089B063 | contains 00000000089A063 | contains 8A0000000899121 |
| pfn 89c | ---DA--KWEV | pfn 89b | ---DA--KWEV |
| 0: kd> !pte poi(nt!MmWriteableSharedUserData) | | | |
| VA fffffd0010402000 | | | |
| PXE at FFFFD6E371B8DF0 | PPE at FFFFD6E371BE000 | PDE at FFFFD6E37C00410 | PTE at FFFFD6F80082010 |
| contains 0A00000010206863 | contains 0A00000010207863 | contains 0A00000010208863 | contains 8A0000000899963 |
| pfn 10206 | ---DA--KWEV | pfn 10207 | ---DA--KWEV |

“Nickel and Dime” Mitigations

- Restricted Caller Improvements*
 - **ExIsRestrictedCaller** was introduced in Windows 8.1 to, for example, block low integrity processes from invoking sensitive APIs that leak kernel addresses by design

* Not the official name of the mitigation

“Nickel and Dime” Mitigations

- **ExIsRestrictedCaller** now uses **SeDebugPrivilege** to determine if kernel addresses can be returned to any caller – not just low integrity processes
 - Effectively this means only admins should be allowed to leak kernel addresses (Microsoft doesn't recognize an admin-to-kernel boundary)

```
bool __fastcall ExIsRestrictedCaller(_KPROCESSOR_MODE PreviousMode, bool *DisallowKernelAddresses)
{
    bool v4; // b1
    _SECURITY_SUBJECT_CONTEXT secSubjectContext; // [rsp+50h] [rbp-28h] BYREF
    int status; // [rsp+80h] [rbp+8h] BYREF
    int grantedAccessMask; // [rsp+88h] [rbp+10h] BYREF

    status = 0;
    grantedAccessMask = 0;
    memset(&secSubjectContext, 0, sizeof(secSubjectContext));
    if ( DisallowKernelAddresses )
        *DisallowKernelAddresses = 0;
    if ( PreviousMode == KernelMode )
        return 0;
    if ( DisallowKernelAddresses )
    {
        Feature_RestrictKernelAddressLeaks__private_IsEnabledPreCheck();
        *DisallowKernelAddresses = !SeSinglePrivilegeCheck(SeDebugPrivilege, PreviousMode);
    }
}
```

“Nickel and Dime” Mitigations

- How does one get kernel addresses from user mode?
 - `ntdll!NtQuerySystemInformation` -> `nt!NtQuerySystemInformation` -> `nt!ExpQuerySystemInformation` is the code path taken to leak kernel object addresses to user-mode via the `SystemHandleInformation` class
 - Even if `nt!ExpQuerySystemInformation` doesn't identify a low integrity process, `ExIsRestrictedCaller` can still override if kernel addresses are allowed
 - Meaning we can restrict even “normal” medium-integrity processes

```
IsRestrictedCaller = ExIsRestrictedCaller(PreviousMode, &disallowKernelAddresses);
if ( !*&IsRestrictedCaller )
{
    SystemBasicInformation = ExpGetHandleInformation(disallowKernelAddresses, userBuffer, v158, &v157);
    EtwTiLogSyscallUsage(SystemBasicInformation, 4u);
    goto LABEL_40;
}
return 0xC0000022i64; // STATUS_ACCESS_DENIED
```


“Nickel and Dime” Mitigations

- How are addresses removed?
 - Example: **ObpCaptureHandleInformation**
 - Eventually called when a caller requests all system handle information (usually to leak an **EPROCESS** or **KTHREAD** object in terms of exploitation)

Default kernel address returned in 0 (NULL)

Only if kernel addresses are allowed are kernel addresses returned to the caller

```
__int64 __fastcall ObpCaptureHandleInformation(
    bool DisallowKernelAddresses,
    _SYSTEM_HANDLE_TABLE_ENTRY_INFO **HandleEntryInfoBuffer,
    __int16 ProcessId,
    void *HandleTableEntry,
    unsigned __int16 HandleIndex,
    unsigned int Length,
    unsigned int *RequiredLength)
{
    unsigned int v11; // eax
    unsigned int status; // r9d
    unsigned __int64 HandlePointer; // r10
    void *actualKernelObjectAddress; // rdx
    char handleTableEntry; // r8

    v11 = *RequiredLength + 24;
    *RequiredLength = v11;
    if ( v11 >= 0x18 )
    {
        if ( Length >= v11 )
        {
            HandlePointer = ExGetHandlePointer(HandleTableEntry);
            LOWORD((*HandleEntryInfoBuffer->ProcessId) = ProcessId;
            status = 0;
            actualKernelObjectAddress = 0i64;
            handleTableEntry = (*HandleTableEntry >> 17) & 7 | 8;
            if ( (*HandleTableEntry + 2) & 0x20000000 == 0 )
            {
                handleTableEntry = (*HandleTableEntry >> 17) & 7;
                (*HandleEntryInfoBuffer->Flags = handleTableEntry & 7;
                (*HandleEntryInfoBuffer->ObjectTypeName = *(ObTypeIndexTable[ObHeaderCookie ^ *(HandlePointer + 24) ^ BYTE1(HandlePointer)]
                    + 40);
                (*HandleEntryInfoBuffer->Handle = HandleIndex;
            }
            if ( !DisallowKernelAddresses )
            {
                actualKernelObjectAddress = (HandlePointer + 0x30); // Compensate for the OBJECT_HEADER (0x30 bytes)
                (*HandleEntryInfoBuffer->Object = actualKernelObjectAddress; // This is NULL if kernel address leaks are not allowed
                HIWORD((*HandleEntryInfoBuffer->ProcessId) = 0;
                (*HandleEntryInfoBuffer)[1].ProcessId = *(HandleTableEntry + 2) & 0x1FFFFFFF;
                *HandleEntryInfoBuffer = (*HandleEntryInfoBuffer + 24);
            }
        }
    }
}
```

“Nickel and Dime” Mitigations

- Why should we care?
 - Previously, adversaries could use **NtQuerySystemInformation** from a medium integrity processes (which are what processes spawn as by default) to easily grab specific **EPROCESS** or **KTHREAD** objects
 - With the new improvements, you must have **SeDebugPrivilege** – which effectively means you need to be admin to leak kernel addresses
 - Admin -> kernel is not a security boundary! If you have admin, you have achieved privilege escalation already

“Nickel and Dime” Mitigations

- **PreviousMode** Corruption Mitigation
 - Each thread object managed by the kernel (KTHREAD) has a member called **PreviousMode**
 - Value comes from the **KPROCESSOR_MODE** enum
 - **KernelMode** = 0
 - **UserMode** = 1
 - **KernelMode**
 - Caller originated from a kernel thread
 - **UserMode**
 - Caller originated from a user thread

C++

```
typedef enum _KPROCESSOR_MODE {  
    KernelMode,  
    UserMode,  
    MaximumMode  
} KPROCESSOR_MODE;
```

“Nickel and Dime” Mitigations

- Why is this important?
 - **PreviousMode** is used in a variety of important security checks when the OS does things like process user-input from a system call, requests/provides a handle, etc.
- Example: **NtTerminateProcess**
 - Process object lookup involves **PreviousMode**

```
if ( processHandle )
{
    result = ObpReferenceObjectByHandleWithTag(
        processHandle,
        1,
        PsProcessType,
        CurrentThread->PreviousMode,
        0x65547350,
        &processObject,
        0i64,
        0i64);
}
```


“Nickel and Dime” Mitigations

- (Example cont'd) Handles are stored in a per-process handle table. If **PreviousMode == KernelMode**, the lookup happens in the System process handle table
 - To prevent abuse, the OS checks if the process handle resides within the “kernel handle range” and the calling thread’s **PreviousMode**
 - If a user-mode thread (**PreviousMode == UserMode**) is detected, it means a user-mode client is attempting to use a kernel-level handle (for example, a privileged object which a kernel-mode component requested) and the OS returns access denied to the user-mode caller

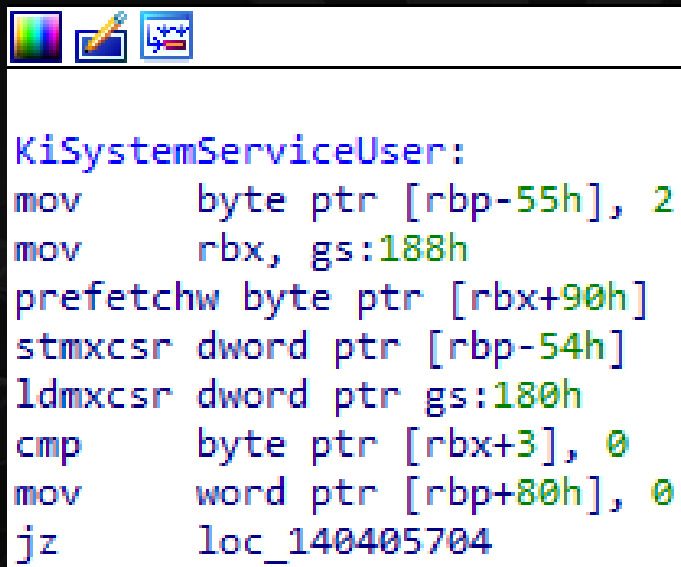
```
if ( ObjectType == PsThreadType || !ObjectType )
{
    if ( (ProcessHandle & 0xFFE00000) != 0 && PreviousMode )
    {
        return 0xC0000022; STATUS_ACCESS_DENIED
    }
}
```

“Nickel and Dime” Mitigations

- With a kernel primitive such as an arbitrary decrement, or arbitrary write, it is possible to “trick” the OS into thinking a user-mode thread is a kernel-mode caller
 1. Prior to invoking a system call, locate the KTHREAD object associated with the thread the system call will be invoked on
 2. Decrement **KTHREAD->PreviousMode** to 0 (**KernelMode**)
 3. Invoke the system call
 4. The thread is now a “kernel-mode caller” and can access privileged objects, write/read to kernel memory, etc.

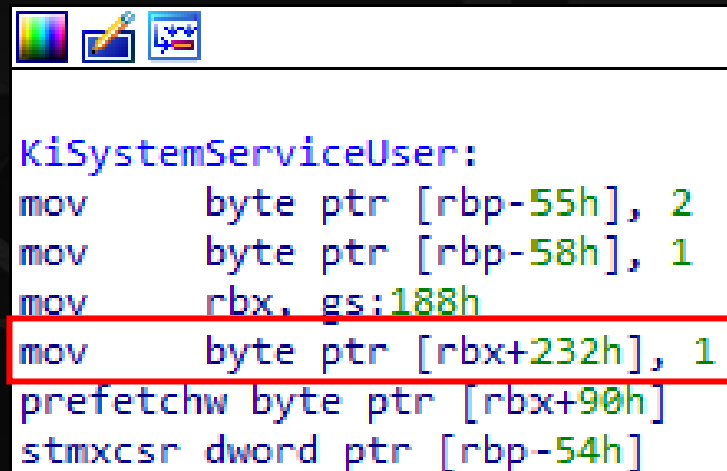
“Nickel and Dime” Mitigations

- With the latest Insider Preview builds, this is no longer the case
 - Firstly, as part of the system call dispatching logic, **PreviousMode** is explicitly set to **UserMode**, from the kernel



```
KiSystemServiceUser:
mov     byte ptr [rbp-55h], 2
mov     rbx, gs:188h
prefetchw byte ptr [rbx+90h]
stmxcsr dword ptr [rbp-54h]
ldmxcsr dword ptr gs:180h
cmp     byte ptr [rbx+3], 0
mov     word ptr [rbp+80h], 0
jz      loc_140405704
```

KTHREAD is fetched into RBX only (pre-mitigation)

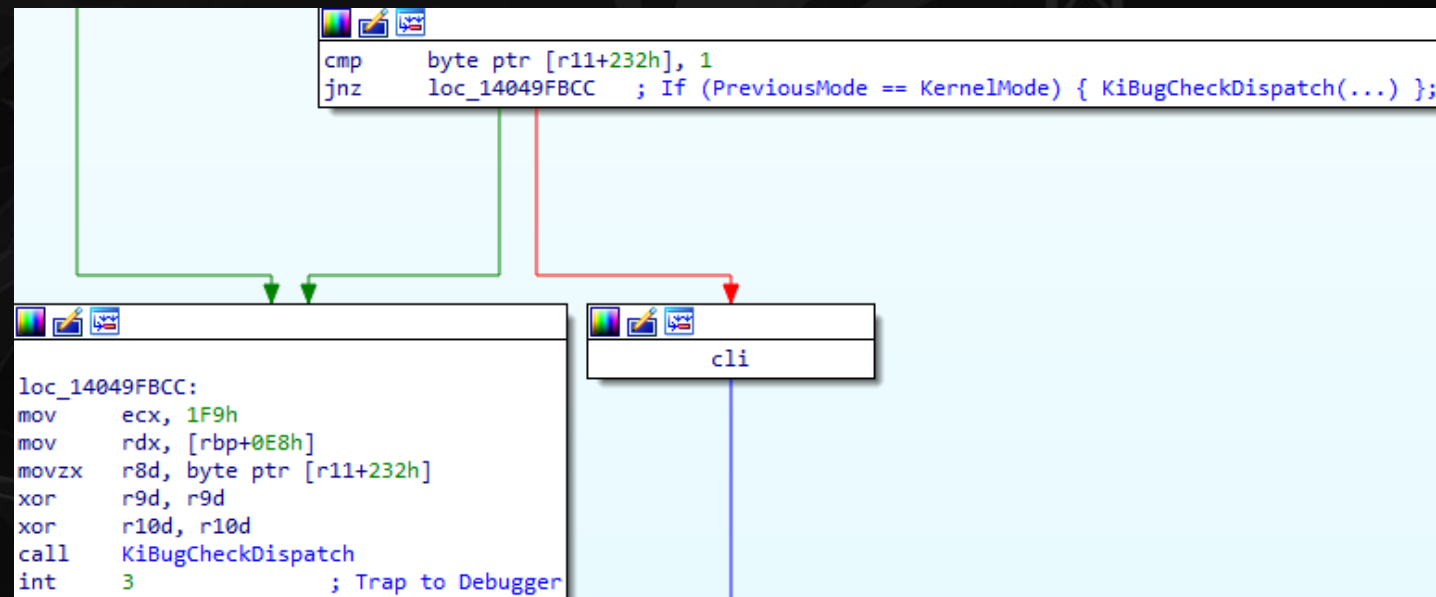


```
KiSystemServiceUser:
mov     byte ptr [rbp-55h], 2
mov     byte ptr [rbp-58h], 1
mov     rbx, gs:188h
mov     byte ptr [rbx+232h], 1
prefetchw byte ptr [rbx+90h]
stmxcsr dword ptr [rbp-54h]
```

KTHREAD is fetched and **PreviousMode** (offset 0x232) is explicitly set to **UserMode** after the system call is invoked (mitigation enabled)

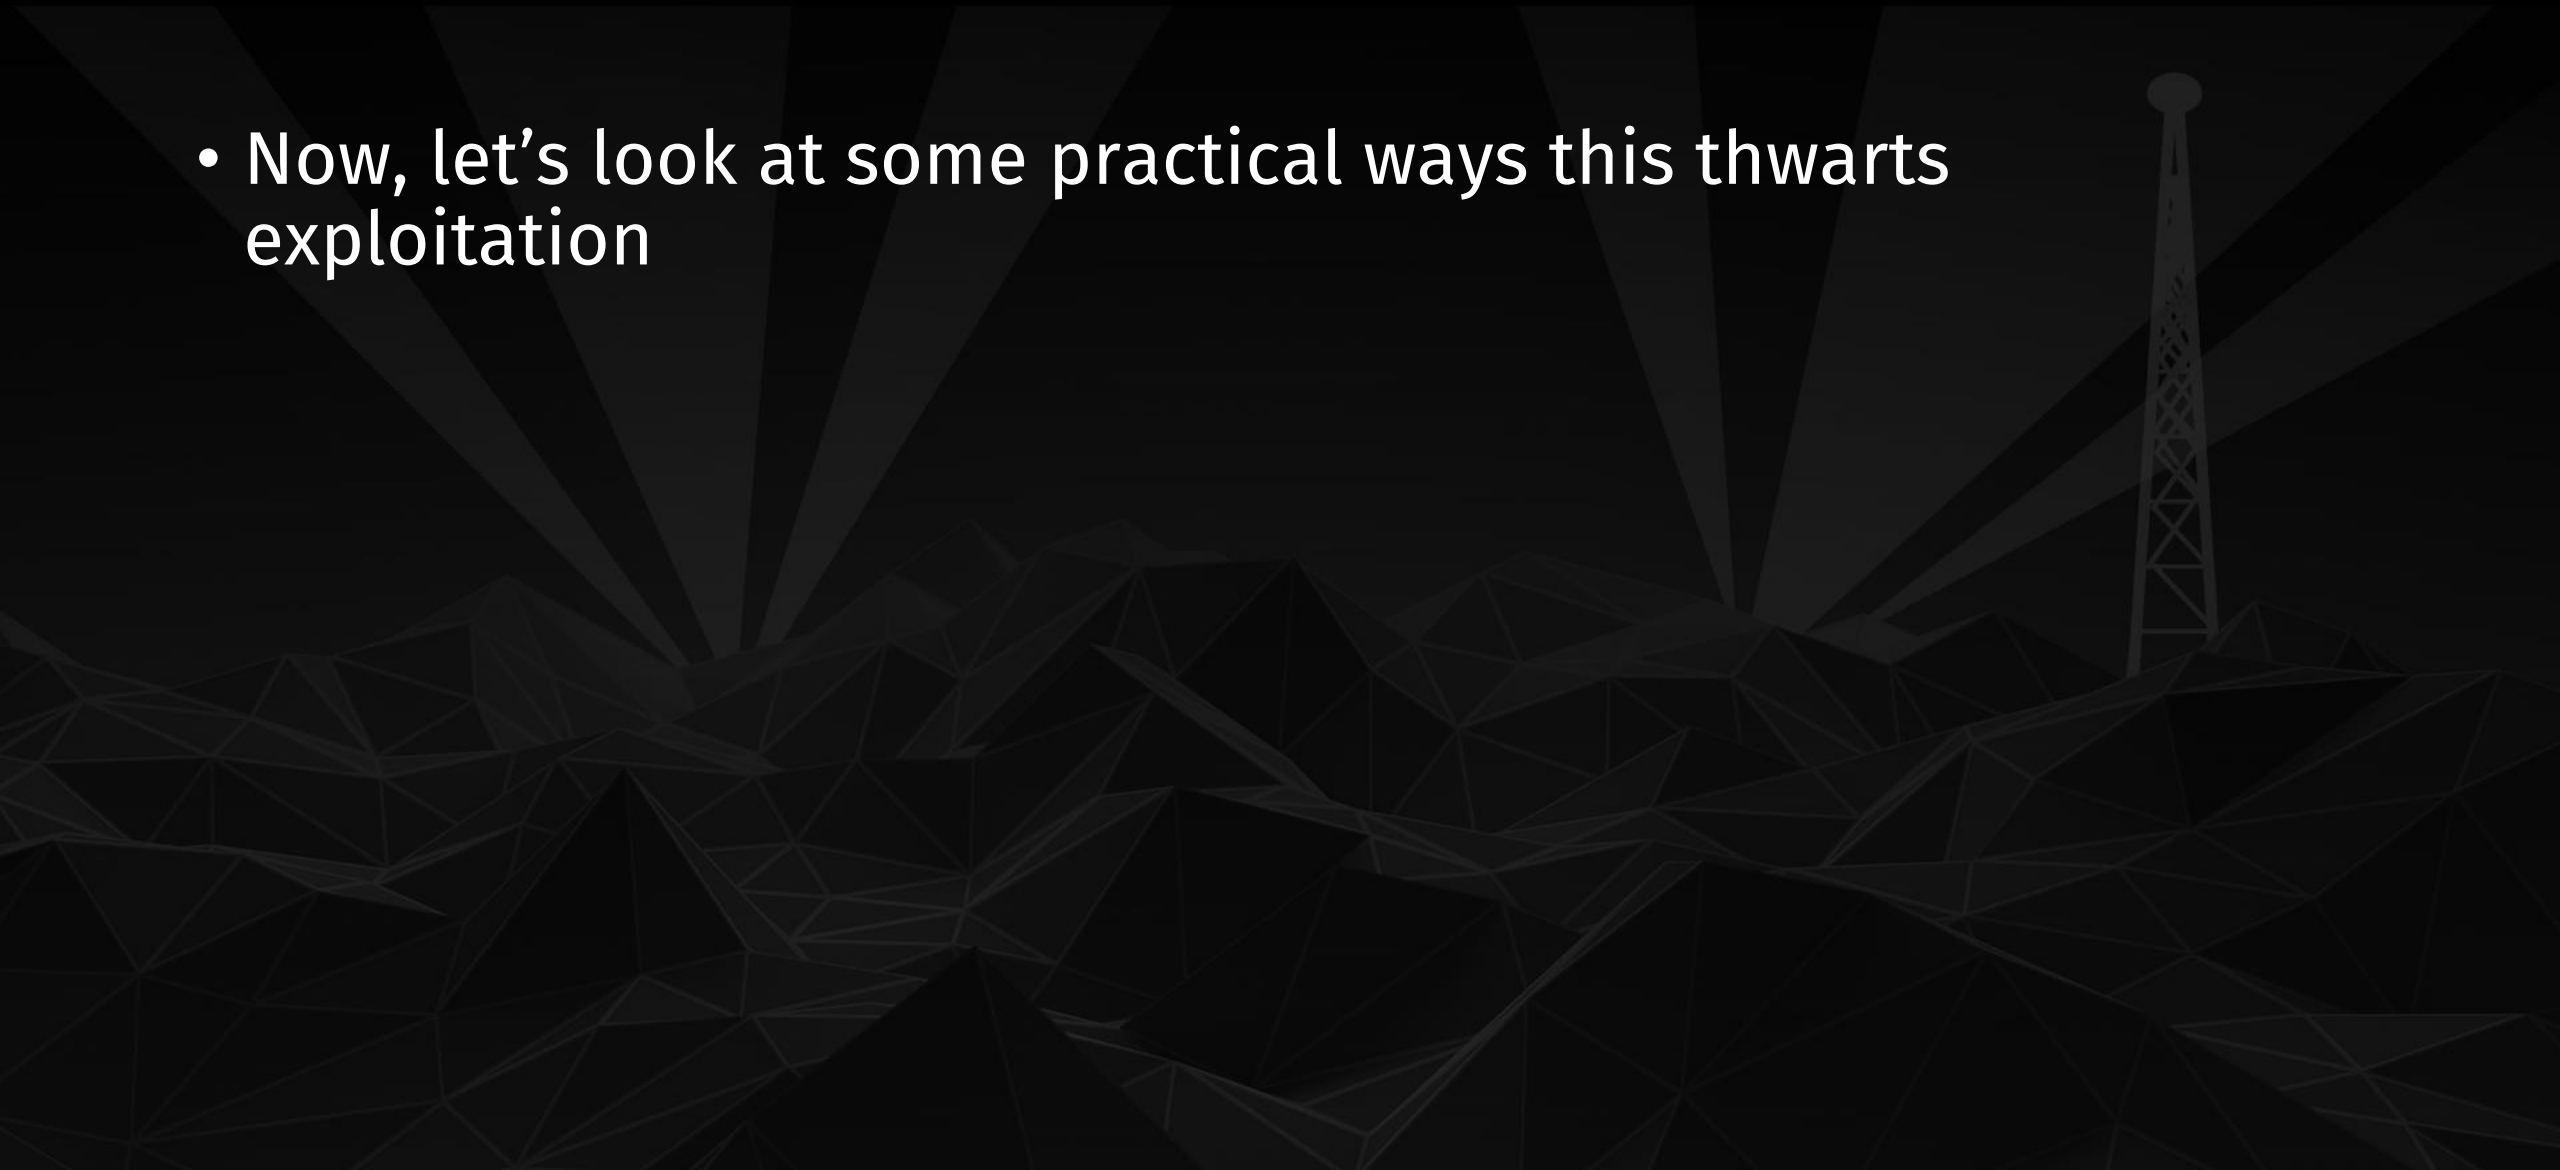
“Nickel and Dime” Mitigations

- In addition, as part of the **sysret** logic, **PreviousMode** is verified again
 - **PreviousMode** cannot be corrupted at the time of the system call invocation, and if it is, it must be restored before the **sysret** – effectively mitigating this technique



“Nickel and Dime” Mitigations

- Now, let's look at some practical ways this thwarts exploitation



Demo #1

- In order to corrupt **PreviousMode**, we need a **KTHREAD** object
 - To get a **KTHREAD** object, without an already existing KASLR/arbitrary read primitive we would normally use **NtQuerySystemInformation**

```
do
{
    //
    // Free the last allocation
    //
    free(handleInformation);

    //
    // Increment the size.
    //
    size = size * 2;

    //
    // Re-alloc.
    //
    handleInformation = reinterpret_cast<PSYSTEM_HANDLE_INFORMATION>(malloc(size));

    if (handleInformation == nullptr)
    {
        goto Exit;
    }

    //
    // This may or may not be successful.
    //
    status = NtQuerySystemInformation(SystemHandleInformation,
                                     handleInformation,
                                     size,
                                     &outSize);
} while (status == STATUS_INFO_LENGTH_MISMATCH);
```

```
if (!NT_SUCCESS(status))
{
    wprintf(L"[-] Error in NtQuerySystemInformation call (0x%lx)!\\n", status);
    goto Exit;
}

for (ULONG i = 0; i < handleInformation->NumberOfHandles; i++)
{
    if ((handleInformation->Handles[i].ProcessId != GetCurrentProcessId()) ||
        (handleInformation->Handles[i].Handle != reinterpret_cast<USHORT>(TargetThread)))
    {
        continue;
    }

    //
    // We found our target thread.
    //
    *ThreadObject = reinterpret_cast<ULONG64>(handleInformation->Handles[i].Object);

    result = true;
    break;
}
```


Demo #2

- Let's say you have a memory disclosure vulnerability, or you have some other primitive and you still manage to leak a **KTHREAD** object, as an example
 - Attempting to corrupt **PreviousMode** then invoke a “privileged system call” is effectively impossible

Conclusion

- Not every effective mitigation has to be a “huge”
- You may not even be aware of a given “nickel/dime” mitigation exists until an exploit starts to break on a new version of Windows
 - Take advantage of Insider Preview, Insider Preview SDK, etc. to look for potential upcoming changes and to look for trends!
- Even smaller mitigations have the potential, depending on your primitives as always, in significantly increase exploitation cost
- Happy exploit writing!

References

- <https://www.alex-ionescu.com/kaslr-bypass-mitigations-in-windows-8-1/>
- https://twitter.com/yarden_shafir/status/1644766239137357829
- <https://twitter.com/GabrielLandau/status/1597001955909697536/>
- <https://github.com/hacksystem/HackSysExtremeVulnerableDriver>